

# kafka传数据到Flink存储到mysql之Flink使用SQL语句聚合数据流（设置时间窗口，EventTime）...

网上没什么资料，就分享下:)

简单模式: kafka传数据到Flink存储到mysql 可以参考网站:

## 利用Flink stream从kafka中写数据到mysql

## maven依赖情况:

```
<project xmlns="http://maven.apache.org/POM/4.0.0"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xsi:schemaLocation="http://maven.apache.org/POM/4.0.0
http://maven.apache.org/xsd/maven-4.0.0.xsd">
<modelVersion>4.0.0</modelVersion>

<groupId>com.xxr</groupId>
<artifactId>flink</artifactId>
<version>0.0.1-SNAPSHOT</version>
<packaging>jar</packaging>
<properties>
<project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
<flink.version>1.4.1</flink.version>
</properties>
<build>
<pluginManagement>
<plugins>
<!-- 设置java版本为1.8 -->
<plugin>
<groupId>org.apache.maven.plugins</groupId>
<artifactId>maven-compiler-plugin</artifactId>
<configuration>
<source>1.8</source>
<target>1.8</target>
<encoding>UTF-8</encoding>
<compilerArgs>
<arg>-extdirs</arg>
```

```

        <arg>${project.basedir}/src/lib</arg>
    </compilerArgs>
</configuration>
</plugin>
<!-- maven-assembly方式打包成jar -->
<plugin>
    <groupId>org.apache.maven.plugins</groupId>
    <artifactId>maven-assembly-plugin</artifactId>
    <version>2.5.5</version>
    <configuration>
        <archive>
            <manifest>

<mainClass>com.xxr.flink.stream_sql</mainClass>
            </manifest>
        </archive>
        <descriptorRefs>
            <descriptorRef>jar-with-
dependencies</descriptorRef>
        </descriptorRefs>
    </configuration>
</plugin>
</plugins>
</pluginManagement>
</build>
<dependencies>
    <dependency>
        <groupId>junit</groupId>
        <artifactId>junit</artifactId>
        <version>3.8.1</version>
        <scope>test</scope>
    </dependency>
    <dependency>
        <groupId>org.slf4j</groupId>
        <artifactId>slf4j-log4j12</artifactId>
        <version>1.8.0-beta1</version>
    </dependency>
    <dependency>
        <groupId>org.slf4j</groupId>
        <artifactId>slf4j-api</artifactId>
        <version>1.8.0-beta1</version>
    </dependency>
    <dependency>
        <groupId>log4j</groupId>
        <artifactId>log4j</artifactId>

```

```

        <version>1.2.17</version>
    </dependency>
    <dependency>
        <groupId>org.scala-lang</groupId>
        <artifactId>scala-compiler</artifactId>
        <version>2.11.1</version>
    </dependency>
    <dependency>
        <groupId>org.scala-lang.modules</groupId>
        <artifactId>scala-xml_2.11</artifactId>
        <version>1.0.2</version>
    </dependency>

    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-java</artifactId>
        <version>${flink.version}</version>
    </dependency>

    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-streaming-java_2.11</artifactId>
        <version>${flink.version}</version>
        <scope>provided</scope>
    </dependency>
    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-streaming-scala_2.11</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-clients_2.11</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <!-- https://mvnrepository.com/artifact/org.apache.flink/flink-core
-->

    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-core</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <!-- https://mvnrepository.com/artifact/org.apache.flink/flink-
runtime -->
    <dependency>

```

```

        <groupId>org.apache.flink</groupId>
        <artifactId>flink-runtime_2.11</artifactId>
        <version>${flink.version}</version>
    </dependency>

    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-connector-wikiedits_2.11</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-connector-kafka-0.8_2.11</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-table_2.11</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <dependency>
        <groupId>org.apache.flink</groupId>
        <artifactId>flink-jdbc</artifactId>
        <version>${flink.version}</version>
    </dependency>
    <dependency>
        <groupId>mysql</groupId>
        <artifactId>mysql-connector-java</artifactId>
        <version>5.1.39</version>
    </dependency>
</dependencies>
</project>

```

配置文件及sql语句，时间窗口是1分钟：

```

public class JDBCTestBase {
    //每过一分钟计算一分钟内的num最大值，以rowtime作为时间基准
    public static final String SQL_MAX = "SELECT  MAX(num)
,TUMBLE_END(rowtime, INTERVAL '1' minute) as wEnd FROM wiki_table  group by
TUMBLE(rowtime, interval '1' minute)";
    public static final String SQL_AVG = "SELECT  AVG(num)
,TUMBLE_END(rowtime, INTERVAL '1' minute) as wEnd FROM wiki_table  group by
TUMBLE(rowtime, interval '1' minute)";
}

```

```

    public static final String SQL_MIN = "SELECT  MIN(num)
,TUMBLE_END(rowtime, INTERVAL '1' minute) as wEnd FROM wiki_table  group by
TUMBLE(rowtime, interval '1' minute)";

    public static final String kafka_group = "test-consumer-group";
    public static final String kafka_zookper = "localhost:2181";
    public static final String kafka_hosts = "localhost:9092";
    public static final String kafka_topic = "wiki-result";
    public static final String DRIVER_CLASS = "com.mysql.jdbc.Driver";
    public static final String DB_URL = "jdbc:mysql://localhost:3306/db?
user=user&password=password";
}

```

## MySQL建表:

```

CREATE TABLE wiki (
  Id int(11) NOT NULL AUTO_INCREMENT,
  avg double DEFAULT NULL,
  time timestamp NULL DEFAULT NULL,
  PRIMARY KEY (Id)
)

```

## 发送数据到kafka,用的是flink example的wikiproducer:

### Monitoring the Wikipedia Edit Stream

```

import org.apache.flink.api.java.functions.KeySelector;
import org.apache.flink.api.java.tuple.Tuple2;
import org.apache.flink.api.java.tuple.Tuple3;
import org.apache.flink.streaming.api.datastream.DataStream;
import org.apache.flink.streaming.api.datastream.KeyedStream;
import
org.apache.flink.streaming.api.environment.StreamExecutionEnvironment;
import org.apache.flink.streaming.api.windowing.time.Time;
import org.apache.flink.streaming.connectors.kafka.FlinkKafkaProducer08;
import org.apache.flink.streaming.connectors.wikiedits.WikipediaEditEvent;
import org.apache.flink.streaming.connectors.wikiedits.WikipediaEditsSource;
import org.apache.flink.api.common.serialization.SimpleStringSchema;
import org.apache.flink.api.common.functions.FoldFunction;
import org.apache.flink.api.common.functions.MapFunction;

```

```

public class WikipediaAnalysis {
    public static void main(String[] args) throws Exception {

        StreamExecutionEnvironment see =
StreamExecutionEnvironment.getExecutionEnvironment();

        DataStream<WikipediaEditEvent> edits = see.addSource(new
WikipediaEditsSource());

        KeyedStream<WikipediaEditEvent, String> keyedEdits = edits
            .keyBy(new KeySelector<WikipediaEditEvent, String>() {
                @Override
                public String getKey(WikipediaEditEvent event) {
                    return event.getUser();
                }
            });

        DataStream<Tuple3<String, Long, Long>> result = keyedEdits
            .timeWindow(Time.seconds(10))
            .fold(new Tuple3<>("", 0L, 0L), new
FoldFunction<WikipediaEditEvent, Tuple3<String, Long, Long>>() {
                @Override
                public Tuple3<String, Long, Long> fold(Tuple3<String, Long, Long>
acc, WikipediaEditEvent event) {
                    acc.f0 = event.getUser().trim();
                    acc.f1 += event.getByteDiff();
                    acc.f2 = System.currentTimeMillis();
                    return acc;
                }
            });

        result
            .map(new MapFunction<Tuple3<String, Long, Long>, String>() {
                @Override
                public String map(Tuple3<String, Long, Long> tuple) {
                    return tuple.toString();
                }
            })
            .addSink(new FlinkKafkaProducer08<>("localhost:9092", "wiki-result",
new SimpleStringSchema()));
        result.print();
        see.execute();
    }
}

```

## 重写RichSinkFunction, 用于写入到mysql中:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.PreparedStatement;
import java.sql.Timestamp;

import org.apache.flink.api.java.tuple.Tuple2;
import org.apache.flink.api.java.tuple.Tuple3;
import org.apache.flink.streaming.api.functions.sink.RichSinkFunction;

import kafka.common.Config;

public class WikiSQLSink extends RichSinkFunction<Tuple3<String, Long, Long>>
{

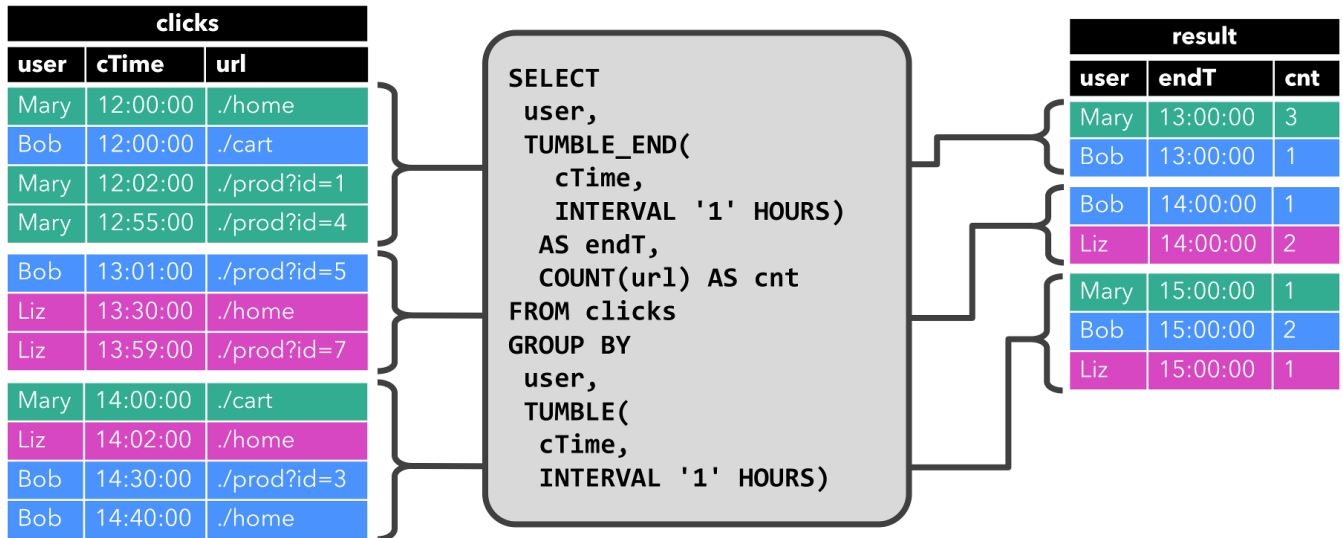
    private static final long serialVersionUID = 1L;
    private Connection connection;
    private PreparedStatement preparedStatement;
    String drivename = JDBCTestBase.DRIVER_CLASS;
    String dburl = JDBCTestBase.DB_URL;

    @Override
    public void invoke(Tuple3<String, Long, Long> value) throws Exception {
        Class.forName(drivename);
        connection = DriverManager.getConnection(dburl);
        String sql = "INSERT into wiki(name,avg,time) values(?,?,?)";
        preparedStatement = connection.prepareStatement(sql);
        preparedStatement.setString(1, value.f0);
        preparedStatement.setLong(2, value.f1);
        preparedStatement.setLong(3, value.f2);
        //preparedStatement.setTimestamp(4, new
Timestamp(System.currentTimeMillis()));
        preparedStatement.executeUpdate();
        if (preparedStatement != null) {
            preparedStatement.close();
        }
        if (connection != null) {
            connection.close();
        }
    }
}
```

```
}
```

```
}
```

用Flink中流计算类，用的是EventTime，用sql语句对数据进行聚合，写入数据到mysql中去，sql的语法用的是另一个开源框架[Apache Cassandra](#)：



图片说明：

<https://ci.apache.org/projects/flink/flink-docs-release-1.4/dev/table/streaming.html#time-attributes>

```
package com.xxr.flink;

import java.sql.Timestamp;
import java.util.Date;
import java.util.Properties;
import java.util.concurrent.TimeUnit;

import org.apache.commons.lang3.StringUtils;
import org.apache.flink.api.common.functions.FilterFunction;
import org.apache.flink.api.common.functions.MapFunction;
import org.apache.flink.api.common.restartstrategy.RestartStrategies;
import org.apache.flink.api.common.typeinfo.TypeInformation;
import org.apache.flink.api.common.typeinfo.Types;
import org.apache.flink.api.java.io.jdbc.JDBCAppendTableSink;
import org.apache.flink.api.java.tuple.Tuple3;
import org.apache.flink.api.java.tuple.Tuple4;
import org.apache.flink.core.fs.FileSystem.WriteMode;
import org.apache.flink.streaming.api.TimeCharacteristic;
```



```

import org.apache.flink.streaming.api.datastream.DataStream;
import
org.apache.flink.streaming.api.environment.StreamExecutionEnvironment;
import org.apache.flink.streaming.api.functions.TimestampAssigner;
import org.apache.flink.streaming.api.windowing.time.Time;
import org.apache.flink.streaming.connectors.kafka.FlinkKafkaConsumer08;
import org.apache.flink.streaming.util.serialization.SimpleStringSchema;
import org.apache.flink.table.api.Table;
import org.apache.flink.table.api.TableEnvironment;
import org.apache.flink.table.api.WindowedTable;
import org.apache.flink.table.api.java.StreamTableEnvironment;
//时间参数网址
//https://ci.apache.org/projects/flink/flink-docs-release-1.4/dev/table/streaming.html#event-time
//Concepts & Common API
//https://ci.apache.org/projects/flink/flink-docs-release-1.4/dev/table/common.html#register-a-table
//SQL语法
//https://ci.apache.org/projects/flink/flink-docs-release-1.4/dev/table/sql.html
public class stream_sql {
    public static void main(String[] args) throws Exception {

        Properties pro = new Properties();
        pro.put("bootstrap.servers", JDBCTestBase.kafka_hosts);
        pro.put("zookeeper.connect", JDBCTestBase.kafka_zookper);
        pro.put("group.id", JDBCTestBase.kafka_group);
        StreamExecutionEnvironment env =
StreamExecutionEnvironment.getExecutionEnvironment();
        StreamTableEnvironment tableEnv =
TableEnvironment.getTableEnvironment(env);
        // env.getConfig().disableSysoutLogging(); //设置此可以屏蔽掉日记打印情况

env.getConfig().setRestartStrategy(RestartStrategies.fixedDelayRestart(4,
10000));

        env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);
        env.enableCheckpointing(5000);

        DataStream<String> sourceStream = env
            .addSource(new FlinkKafkaConsumer08<String>
(JDBCTestBase.kafka_topic, new SimpleStringSchema(), pro));

        DataStream<Tuple3<Long, String, Long>> sourceStreamTra =
sourceStream.filter(new FilterFunction<String>() {
            @Override

```

```

        public boolean filter(String value) throws Exception {
            return StringUtils.isNotBlank(value);
        }
    }).map(new MapFunction<String, Tuple3<Long, String, Long>>() {
        @Override
        public Tuple3<Long, String, Long> map(String value) throws
Exception {
            String temp = value.replaceAll("(\\(|\\|\\|\\|)", "");
            String[] args = temp.split(",");
            try {
                return new Tuple3<Long, String, Long>
(Long.valueOf(args[2]), args[0].trim(), Long.valueOf(args[1]));

            } catch (Exception e) {
                // TODO Auto-generated catch block
                e.printStackTrace();
                return new Tuple3<Long, String, Long>
(System.currentTimeMillis(), args[0].trim(), 0L);

            }
        }
    });
    //設置將哪个字段用于eventTime
    DataStream<Tuple3<Long, String, Long>> withTimestampsAndWatermarks =
sourceStreamTra
        .assignTimestampsAndWatermarks(new FirstTandW());
    //内置参数rowtime.rowtime就是EventTime protime是ProcessingTime
    tableEnv.registerDataStream("wiki_table",
withTimestampsAndWatermarks, "etime,name, num,rowtime.rowtime");
    withTimestampsAndWatermarks.print();
    // define sink for room data and execute query
    JDBCAppendTableSink sink =
JDBCAppendTableSink.builder().setDrivername(JDBCTestBase.DRIVER_CLASS)
        .setDBUrl(JDBCTestBase.DB_URL).setQuery("INSERT INTO wiki
(avg,time) VALUES (?,?)")
        .setParameterTypes(Types.LONG, Types.SQL_TIMESTAMP).build();
    //执行查询
    Table result = tableEnv.sqlQuery(JDBCTestBase.SQL_MIN);
    //写入csv
    //      result.writeToSink(new CsvTableSink("D:/a.csv", // output path
    //      "|", // optional: delimit files by '|'
    //      1, // optional: write to a single file
    //      WriteMode.OVERWRITE)); // optional: override existing
files
    //写入数据库

```

```

        result.writeToSink(sink);

        env.execute();

    }
}

```

重写AssignerWithPeriodicWatermarks设置watermark，处理时间是EventTime的话必须要有这个方法，ProcessingTime 可忽略

```

import org.apache.flink.api.java.tuple.Tuple3;
import
org.apache.flink.streaming.api.functions.AssignerWithPeriodicWatermarks;
import org.apache.flink.streaming.api.watermark.Watermark;

public class FirstTandW implements
AssignerWithPeriodicWatermarks<Tuple3<Long,String, Long>> {

    private final long maxOutOfOrderness = 3500; // 3.5 seconds

    private long currentMaxTimestamp;

    @Override
    public long extractTimestamp(Tuple3<Long,String, Long> element, long
previousElementTimestamp) {
        // TODO Auto-generated method stub
        long timestamp = element.f0;
        currentMaxTimestamp = Math.max(timestamp, currentMaxTimestamp);
        return timestamp;
    }

    @Override
    public Watermark getCurrentWatermark() {
        // TODO Auto-generated method stub
        return new Watermark(currentMaxTimestamp - maxOutOfOrderness);
    }

}

```

maven assembly打包成jar，放flink运行就行了，不会打包看我博客

基础知识

## **Flink 的Window 操作**

### **Flink流计算编程--在WindowedStream中体会EventTime与ProcessingTime**

Flink文档写的很好。。刚开始做没仔细看，坑不少

git: <https://github.com/xxrznj/flink-kafka-sql>