

My flink 异常问题

```
1
2 Fail to run sql command: insert overwrite md_v1 select
3     loan_id,
4     sum(posting_amount) as accu_repay_total_amt
5 from
6     ods_txn_tl_payment_allocation_log
7 where
8     account_type = '09'
9 group by
10    loan_id
11 java.io.IOException: org.apache.flink.table.api.ValidationException: SQL validation failed.
    From line 5, column 2 to line 5, column 34: Object 'ods_txn_tl_payment_allocation_log' not
    found
12     at
    org.apache.zepplin.flink.FlinkSqlInterrpeter.callInsertInto(FlinkSqlInterrpeter.java:536)
13     at
    org.apache.zepplin.flink.FlinkSqlInterrpeter.callCommand(FlinkSqlInterrpeter.java:272)
14     at
    org.apache.zepplin.flink.FlinkSqlInterrpeter.runSqlList(FlinkSqlInterrpeter.java:159)
15     at org.apache.zepplin.flink.FlinkSqlInterrpeter.interpret(FlinkSqlInterrpeter.java:124)
16     at
    org.apache.zepplin.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpreter.java:110)
17     at
    org.apache.zepplin.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInt
    erpreterServer.java:776)
18     at
    org.apache.zepplin.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInt
    erpreterServer.java:668)
19     at org.apache.zepplin.scheduler.Job.run(Job.java:172)
20     at org.apache.zepplin.scheduler.AbstractScheduler.runJob(AbstractScheduler.java:130)
21     at
    org.apache.zepplin.scheduler.ParallelScheduler.lambda$runJobInScheduler$0(ParallelScheduler
    .java:39)
22     at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
23     at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
24     at java.lang.Thread.run(Thread.java:748)
25 Caused by: org.apache.flink.table.api.ValidationException: SQL validation failed. From line
    5, column 2 to line 5, column 34: Object 'ods_txn_tl_payment_allocation_log' not found
26     at
    org.apache.flink.table.planner.calcite.FlinkPlannerImpl.org$apache$flink$table$planner$calci
    te$FlinkPlannerImpl$$validate(FlinkPlannerImpl.scala:146)
27     at
    org.apache.flink.table.planner.calcite.FlinkPlannerImpl.validate(FlinkPlannerImpl.scala:108)
28     at
    org.apache.flink.table.planner.operations.SqlToOperationConverter.convert(SqlToOperationConv
```

```
erter.java:185)
29     at
org.apache.flink.table.planner.operations.SqlToOperationConverter.convertSqlInsert(SqlToOperationConverter.java:525)
30     at
org.apache.flink.table.planner.operations.SqlToOperationConverter.convert(SqlToOperationConverter.java:202)
31     at org.apache.flink.table.planner.delegation.ParserImpl.parse(ParserImpl.java:78)
32     at
org.apache.flink.table.api.internal.TableEnvironmentImpl.sqlUpdate(TableEnvironmentImpl.java:753)
33     at
org.apache.zepplin.flink.FlinkSqlInterpreter.callInsertInto(FlinkSqlInterpreter.java:528)
34     ... 12 more
35 Caused by: org.apache.calcite.runtime.CalciteContextException: From line 5, column 2 to line 5, column 34: Object 'ods_txn_tl_payment_allocation_log' not found
36     at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
37     at
sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)
38     at
sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
39     at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
40     at org.apache.calcite.runtime.Resources$ExInstWithCause.ex(Resources.java:457)
41     at org.apache.calcite.sql.SqlUtil.newContextException(SqlUtil.java:839)
42     at org.apache.calcite.sql.SqlUtil.newContextException(SqlUtil.java:824)
43     at
org.apache.calcite.sql.validate.SqlValidatorImpl.newValidationError(SqlValidatorImpl.java:508)
44     at
org.apache.calcite.sql.validate.IdentifierNamespace.resolveImpl(IdentifierNamespace.java:179)
45     at
org.apache.calcite.sql.validate.IdentifierNamespace.validateImpl(IdentifierNamespace.java:184)
46     at org.apache.calcite.sql.validate.AbstractNamespace.validate(AbstractNamespace.java:84)
47     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateNamespace(SqlValidatorImpl.java:1110)
48     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateQuery(SqlValidatorImpl.java:1084)
49     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateFrom(SqlValidatorImpl.java:3256)
50     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateFrom(SqlValidatorImpl.java:3238)
51     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateSelect(SqlValidatorImpl.java:3510)
52     at org.apache.calcite.sql.validate.SelectNamespace.validateImpl(SelectNamespace.java:60)
53     at org.apache.calcite.sql.validate.AbstractNamespace.validate(AbstractNamespace.java:84)
54     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateNamespace(SqlValidatorImpl.java:1110)
55     at
org.apache.calcite.sql.validate.SqlValidatorImpl.validateQuery(SqlValidatorImpl.java:1084)
56     at org.apache.calcite.sql.SqlSelect.validate(SqlSelect.java:232)
57     at
```

```

    org.apache.calcite.sql.validate.SqlValidatorImpl.validateScopedExpression(SqlValidatorImpl.java:1059)
58     at org.apache.calcite.sql.validate.SqlValidatorImpl.validate(SqlValidatorImpl.java:766)
59     at
    org.apache.flink.table.planner.calcite.FlinkPlannerImpl.org$apache$flink$table$planner$calcite$FlinkPlannerImpl$$validate(FlinkPlannerImpl.scala:141)
60     ... 19 more
61 Caused by: org.apache.calcite.sql.validate.SqlValidatorException: Object
    'ods_txn_tl_payment_allocation_log' not found
62     at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
63     at
    sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)
64     at
    sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
65     at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
66     at org.apache.calcite.runtime.Resources$ExInstWithCause.ex(Resources.java:457)
67     at org.apache.calcite.runtime.Resources$ExInst.ex(Resources.java:550)
68     ... 38 more
69
70

```

```

1 Fail to run sql command: insert overwrite md_v1 select
2   loan_id,
3   sum(posting_amount) as accu_repay_total_amt
4 from
5   ods_v1
6 where
7   account_type = '09'
8 group by
9   loan_id
10 java.io.IOException: org.apache.flink.table.api.ValidationException: INSERT OVERWRITE
    requires JDBC:PostgreSQL DynamicTableSink to implement SupportsOverwrite interface.
11     at
    org.apache.zeppelein.flink.FlinkSqlInterrpeter.callInsertInto(FlinkSqlInterrpeter.java:536)
12     at
    org.apache.zeppelein.flink.FlinkSqlInterrpeter.callCommand(FlinkSqlInterrpeter.java:272)
13     at
    org.apache.zeppelein.flink.FlinkSqlInterrpeter.runSqlList(FlinkSqlInterrpeter.java:159)
14     at org.apache.zeppelein.flink.FlinkSqlInterrpeter.interpret(FlinkSqlInterrpeter.java:124)
15     at
    org.apache.zeppelein.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpreter.java:110)
16     at
    org.apache.zeppelein.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInterpreterServer.java:776)
17     at
    org.apache.zeppelein.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInterpreterServer.java:668)
18     at org.apache.zeppelein.scheduler.Job.run(Job.java:172)
19     at org.apache.zeppelein.scheduler.AbstractScheduler.runJob(AbstractScheduler.java:130)
20     at
    org.apache.zeppelein.scheduler.ParallelScheduler.lambda$runJobInScheduler$0(ParallelScheduler.java:39)
21     at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
22     at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)

```

```

23     at java.lang.Thread.run(Thread.java:748)
24 Caused by: org.apache.flink.table.api.ValidationException: INSERT OVERWRITE requires
JDBC:PostgreSQL DynamicTableSink to implement SupportsOverwrite interface.
25     at
org.apache.flink.table.planner.sinks.TableSinkUtils$.validateTableSink(TableSinkUtils.scala:
194)
26     at
org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$2.apply(PlannerBase.scala:227
)
27     at
org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$2.apply(PlannerBase.scala:204
)
28     at scala.Option.map(Option.scala:146)
29     at
org.apache.flink.table.planner.delegation.PlannerBase.translateToRel(PlannerBase.scala:204)
30     at
org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$1.apply(PlannerBase.scala:163
)
31     at
org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$1.apply(PlannerBase.scala:163
)
32     at scala.collection.TraversableLike$$anonfun$map$1.apply(TraversableLike.scala:234)
33     at scala.collection.TraversableLike$$anonfun$map$1.apply(TraversableLike.scala:234)
34     at scala.collection.Iterator$class.foreach(Iterator.scala:891)
35     at scala.collection.AbstractIterator.foreach(Iterator.scala:1334)
36     at scala.collection.IterableLike$class.foreach(IterableLike.scala:72)
37     at scala.collection.AbstractIterable.foreach(Iterable.scala:54)
38     at scala.collection.TraversableLike$class.map(TraversableLike.scala:234)
39     at scala.collection.AbstractTraversable.map(Traversable.scala:104)
40     at
org.apache.flink.table.planner.delegation.PlannerBase.translate(PlannerBase.scala:163)
41     at
org.apache.flink.table.api.internal.TableEnvironmentImpl.translate(TableEnvironmentImpl.java
:1264)
42     at
org.apache.flink.table.api.internal.TableEnvironmentImpl.translateAndClearBuffer(TableEnviro
nmentImpl.java:1256)
43     at
org.apache.flink.table.api.internal.TableEnvironmentImpl.execute(TableEnvironmentImpl.java:1
213)
44     at
org.apache.zepplin.flink.FlinkSqlInterpreter.callInsertInto(FlinkSqlInterpreter.java:530)
45     ... 12 more
46

```

```

1 Fail to run sql command: CREATE TABLE ods_v1 (
2     loan_id varchar,
3     posting_amount decimal,
4     account_type varchar,
5     primary key(loan_id)
6 ) WITH (
7     'connector' = 'jdbc',
8     'url' = 'jdbc:postgresql://172.16.101.37:5432/finance_cloud',
9     'driver' = 'org.postgresql.Driver',
10    'table-name' = 'ods_txn_tl_payment_allocation_log',

```

```

11     'username' = 'gpadmin',
12     'password' = 'gpadmin'
13 )
14 org.apache.flink.table.api.ValidationException: Flink doesn't support ENFORCED mode for
PRIMARY KEY constaint. ENFORCED/NOT ENFORCED controls if the constraint checks are
performed on the incoming/outgoing data. Flink does not own the data therefore the only
supported mode is the NOT ENFORCED mode
15     at
org.apache.flink.table.planner.operations.SqlToOperationConverter.validateTableConstraint(Sql
lToOperationConverter.java:754)
16     at java.util.ArrayList.forEach(ArrayList.java:1257)
17     at
org.apache.flink.table.planner.operations.SqlCreateTableConverter.convertCreateTable(SqlCrea
teTableConverter.java:75)
18     at
org.apache.flink.table.planner.operations.SqlToOperationConverter.convert(SqlToOperationConv
erter.java:188)
19     at org.apache.flink.table.planner.delegation.ParserImpl.parse(ParserImpl.java:78)
20     at
org.apache.flink.table.api.internal.TableEnvironmentImpl.sqlUpdate(TableEnvironmentImpl.java
:753)
21     at
org.apache.zeppelein.flink.FlinkSqlInterrpeter.callCreateTable(FlinkSqlInterrpeter.java:369)
22     at
org.apache.zeppelein.flink.FlinkSqlInterrpeter.callCommand(FlinkSqlInterrpeter.java:275)
23     at
org.apache.zeppelein.flink.FlinkSqlInterrpeter.runSqlList(FlinkSqlInterrpeter.java:159)
24     at org.apache.zeppelein.flink.FlinkSqlInterrpeter.interpret(FlinkSqlInterrpeter.java:124)
25     at
org.apache.zeppelein.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpreter.java:110)
26     at
org.apache.zeppelein.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInt
erpreterServer.java:776)
27     at
org.apache.zeppelein.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInt
erpreterServer.java:668)
28     at org.apache.zeppelein.scheduler.Job.run(Job.java:172)
29     at org.apache.zeppelein.scheduler.AbstractScheduler.runJob(AbstractScheduler.java:130)
30     at
org.apache.zeppelein.scheduler.ParallelScheduler.lambda$runJobInScheduler$0(ParallelScheduler
.java:39)
31     at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
32     at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
33     at java.lang.Thread.run(Thread.java:748)
34

```

```

1 Fail to run sql command: insert overwrite md_v1 select
2     loan_id,
3     sum(posting_amount) as accu_repay_total_amt
4 from
5     ods_v1
6 where
7     account_type = '09'
8 group by
9     loan_id

```

```

10 java.io.IOException: org.apache.flink.table.api.TableException:
   'default_catalog.default_database.md_v1' is a partitioned table, but the underlying
   [JDBC:PostgreSQL] DynamicTableSink doesn't implement SupportsPartitioning interface.
11     at
   org.apache.zepelin.flink.FlinkSqlInterrpeter.callInsertInto(FlinkSqlInterrpeter.java:536)
12     at
   org.apache.zepelin.flink.FlinkSqlInterrpeter.callCommand(FlinkSqlInterrpeter.java:272)
13     at
   org.apache.zepelin.flink.FlinkSqlInterrpeter.runSqlList(FlinkSqlInterrpeter.java:159)
14     at org.apache.zepelin.flink.FlinkSqlInterrpeter.interpret(FlinkSqlInterrpeter.java:124)
15     at
   org.apache.zepelin.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpreter.java:110)
16     at
   org.apache.zepelin.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInt
   erpreterServer.java:776)
17     at
   org.apache.zepelin.interpreter.remote.RemoteInterpreterServer$InterpretJob.jobRun(RemoteInt
   erpreterServer.java:668)
18     at org.apache.zepelin.scheduler.Job.run(Job.java:172)
19     at org.apache.zepelin.scheduler.AbstractScheduler.runJob(AbstractScheduler.java:130)
20     at
   org.apache.zepelin.scheduler.ParallelScheduler.lambda$runJobInScheduler$0(ParallelScheduler
   .java:39)
21     at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
22     at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
23     at java.lang.Thread.run(Thread.java:748)
24 Caused by: org.apache.flink.table.api.TableException:
   'default_catalog.default_database.md_v1' is a partitioned table, but the underlying
   [JDBC:PostgreSQL] DynamicTableSink doesn't implement SupportsPartitioning interface.
25     at
   org.apache.flink.table.planner.sinks.TableSinkUtils$.validateTableSink(TableSinkUtils.scala:
   172)
26     at
   org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$2.apply(PlannerBase.scala:227
   )
27     at
   org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$2.apply(PlannerBase.scala:204
   )
28     at scala.Option.map(Option.scala:146)
29     at
   org.apache.flink.table.planner.delegation.PlannerBase.translateToRel(PlannerBase.scala:204)
30     at
   org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$1.apply(PlannerBase.scala:163
   )
31     at
   org.apache.flink.table.planner.delegation.PlannerBase$$anonfun$1.apply(PlannerBase.scala:163
   )
32     at scala.collection.TraversableLike$$anonfun$map$1.apply(TraversableLike.scala:234)
33     at scala.collection.TraversableLike$$anonfun$map$1.apply(TraversableLike.scala:234)
34     at scala.collection.Iterator$class.foreach(Iterator.scala:891)
35     at scala.collection.AbstractIterator.foreach(Iterator.scala:1334)
36     at scala.collection.IterableLike$class.foreach(IterableLike.scala:72)
37     at scala.collection.AbstractIterable.foreach(Iterable.scala:54)
38     at scala.collection.TraversableLike$class.map(TraversableLike.scala:234)
39     at scala.collection.AbstractTraversable.map(Traversable.scala:104)
40     at

```

```

    org.apache.flink.table.planner.delegation.PlannerBase.translate(PlannerBase.scala:163)
41   at
    org.apache.flink.table.api.internal.TableEnvironmentImpl.translate(TableEnvironmentImpl.java
:1264)
42   at
    org.apache.flink.table.api.internal.TableEnvironmentImpl.translateAndClearBuffer(TableEnviro
nmentImpl.java:1256)
43   at
    org.apache.flink.table.api.internal.TableEnvironmentImpl.execute(TableEnvironmentImpl.java:1
213)
44   at
    org.apache.zeppelein.flink.FlinkSqlInterrpeter.callInsertInto(FlinkSqlInterrpeter.java:530)
45   ... 12 more

```

```

1  Caused by: org.apache.flink.kafka.shaded.org.apache.kafka.common.KafkaException: class
    org.apache.kafka.common.serialization.ByteArrayDeserializer is not an instance of
    org.apache.flink.kafka.shaded.org.apache.kafka.common.serialization.Deserializer
2    at
    org.apache.flink.kafka.shaded.org.apache.kafka.common.config.AbstractConfig.getConfiguredIns
tance(AbstractConfig.java:374)
3    at
    org.apache.flink.kafka.shaded.org.apache.kafka.common.config.AbstractConfig.getConfiguredIns
tance(AbstractConfig.java:392)
4    at org.apache.flink.kafka.shaded.org.apache.kafka.clients.consumer.KafkaConsumer.<init>
    (KafkaConsumer.java:713)
5
6
7

```

```

1
2  org.apache.flink.table.api.ValidationException: Could not find any factory for identifier
    'elasticsearch-7' that implements
    'org.apache.flink.table.factories.DynamicTableSourceFactory' in the classpath.
3

```

4 看起来你直接去读注册的这个 es 表了， es 表是不支持读，只支持写的

```

5
6

```

```

1
2  flink-sql-connector-elasticsearch6_2.11-1.11.1.jar
3  flink-sql-connector-elasticsearch7_2.11-1.11.1.jar
4

```

5 支持不同版本的jar不能一起放在lib下，否则会报各种错误！！

```

6
7

```

@云邪(Jark) 大佬，周末快乐，想问一下我用flink table1.11.1从kafka读取数据，然后创建一张经过agg算子后名为 test 的视图：（x,y 都是String类型）

CREATE VIEW test AS

```
SELECT
    x,
    y,
    cast(count(1) AS varchar) AS ct
FROM kafka
GROUP BY x, y
```

然后从这个视图表grouby然后从这个视图里再进行一些处理：

```
select
    x,
    LISTAGG(CONCAT_WS('=', y, ct), ',') AS lists
FROM test
GROUP BY x
```

从这个视图表grouby x做聚合agg操作，但是会报错，报的是
org.apache.flink.table.data.binary.BinaryRawValueData cannot be cast to
org.apache.flink.table.data.StringData

可是我从视图如果不做视图操作，这样是没有问题：

```
select x, CONCAT_WS('=', y, ct) from test
```

说明count(1) 是能够被cast 成string的，可是经过agg怎么就不行了，请问有什么比较好的办法解决这个问题？

@俞剑波 之前我遇到过，可以写一个自定义强转函数，在sql里面用自己的函数来转换位string